

# GPU Computing



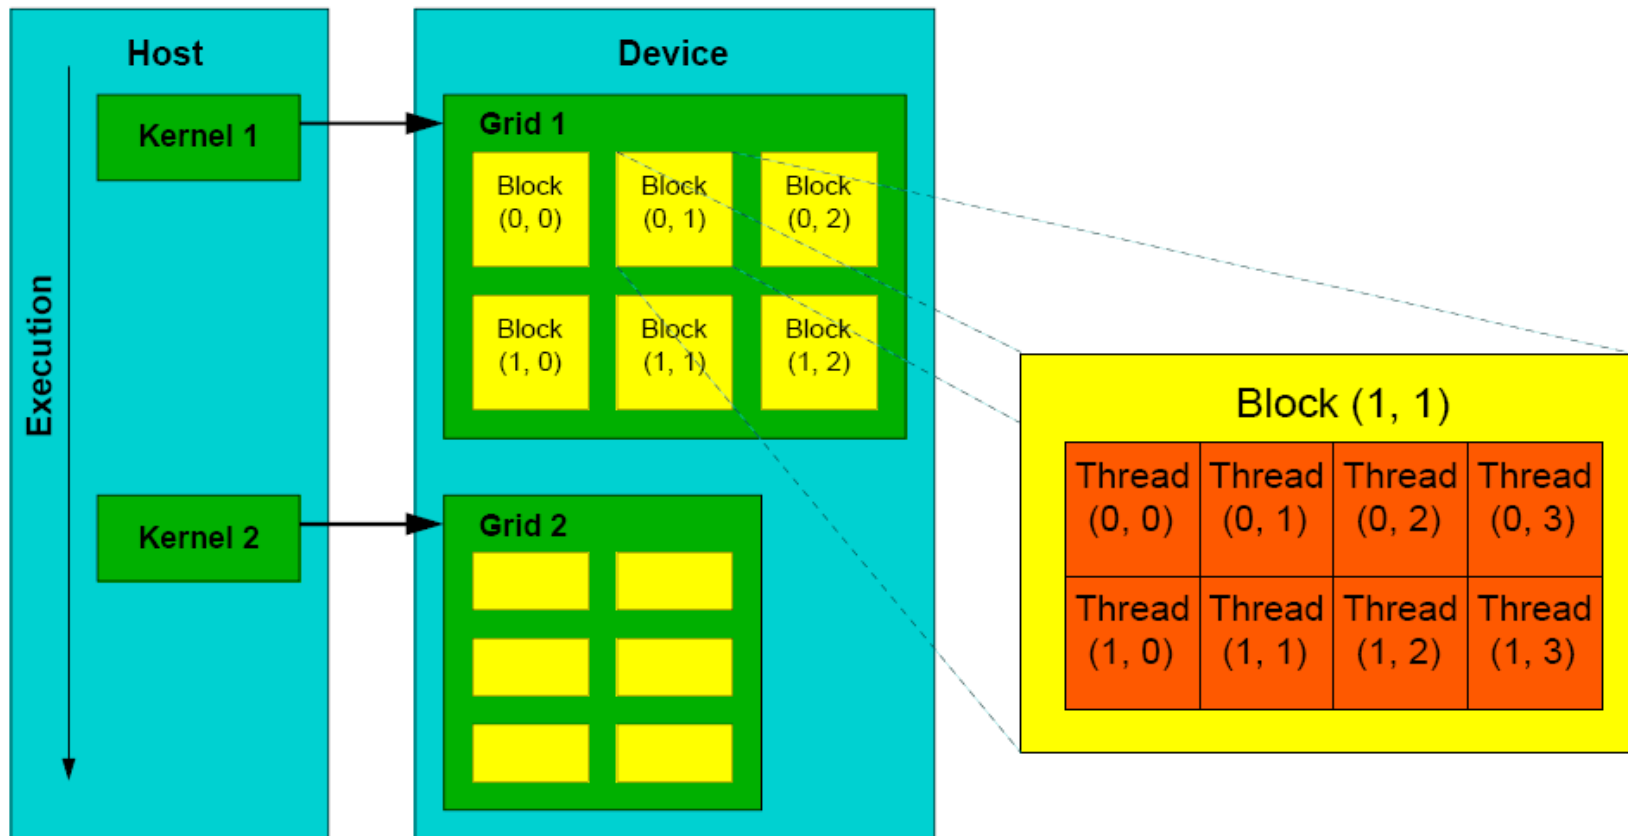


# GPU 硬件架构综述

Hui Liu


Email: [hui.sc.liu@gmail.com](mailto:hui.sc.liu@gmail.com)

- Grids of Blocks
- Blocks of Threads

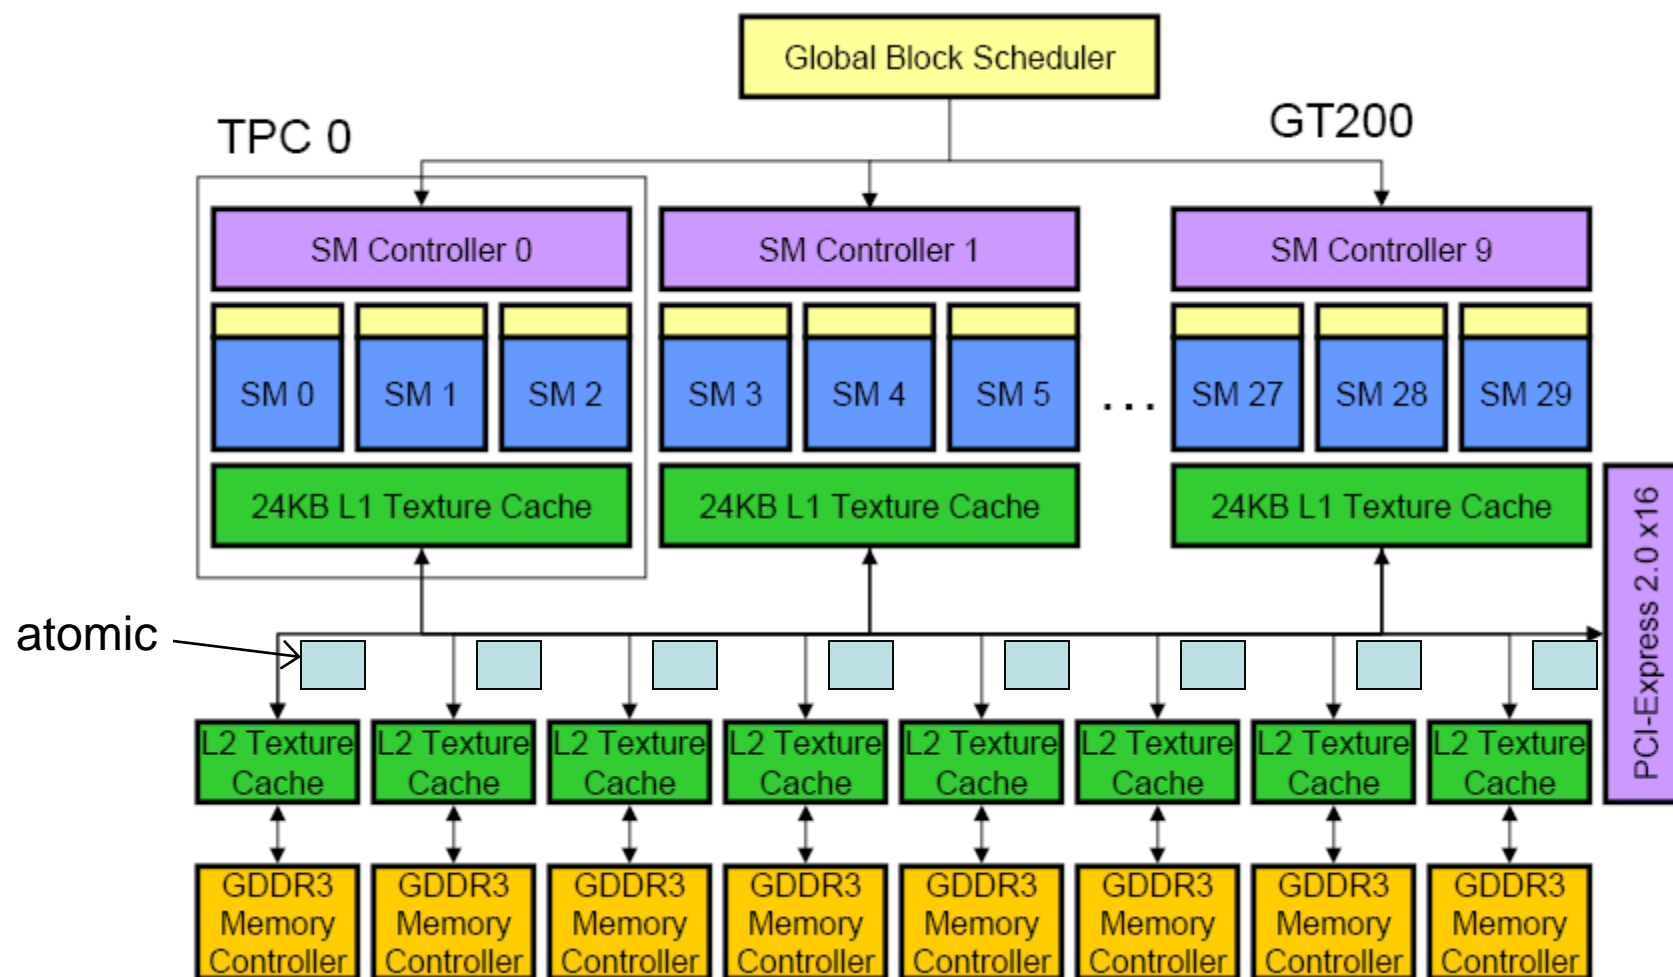


# Thread Blocks Refresher

- Programmer declares (Thread) Block:
  - Block size 1 to **512 (1024, 2048)** concurrent threads
  - Block shape 1D, 2D, or 3D
  - Block dimensions in threads
- All threads in a Block execute the same thread program
- Threads have **thread id** numbers within Block
- Threads share data and synchronize while doing their share of the work
- Thread program uses **thread id** to select work and address shared data



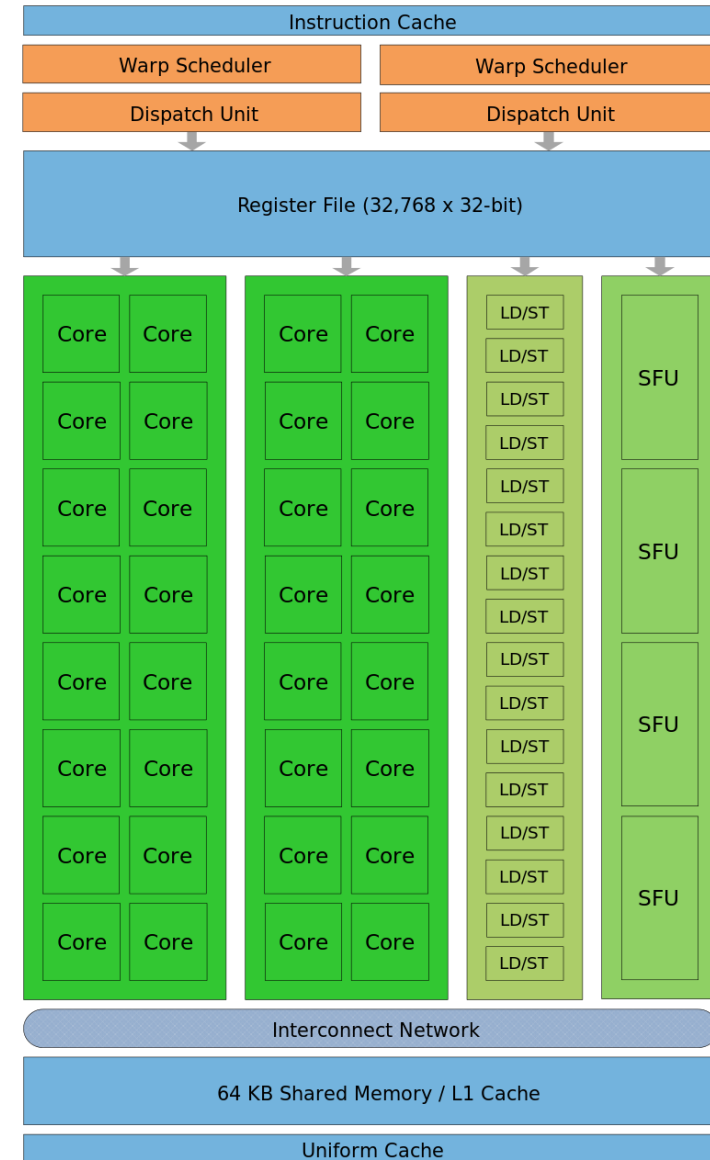
# GT200 Architecture Overview



- **SPA**
  - Streaming Processor Array
- **TPC/GPC**
  - Texture (Graphics) Processor Cluster
    - 3 SM + TEX
- **SM**
  - Streaming Multiprocessor (32 SP)
  - Multi-threaded processor core
  - Fundamental processing unit for CUDA thread block
- **SP (CUDA Core)**
  - Streaming Processor

# Stream Multiprocessor Overview

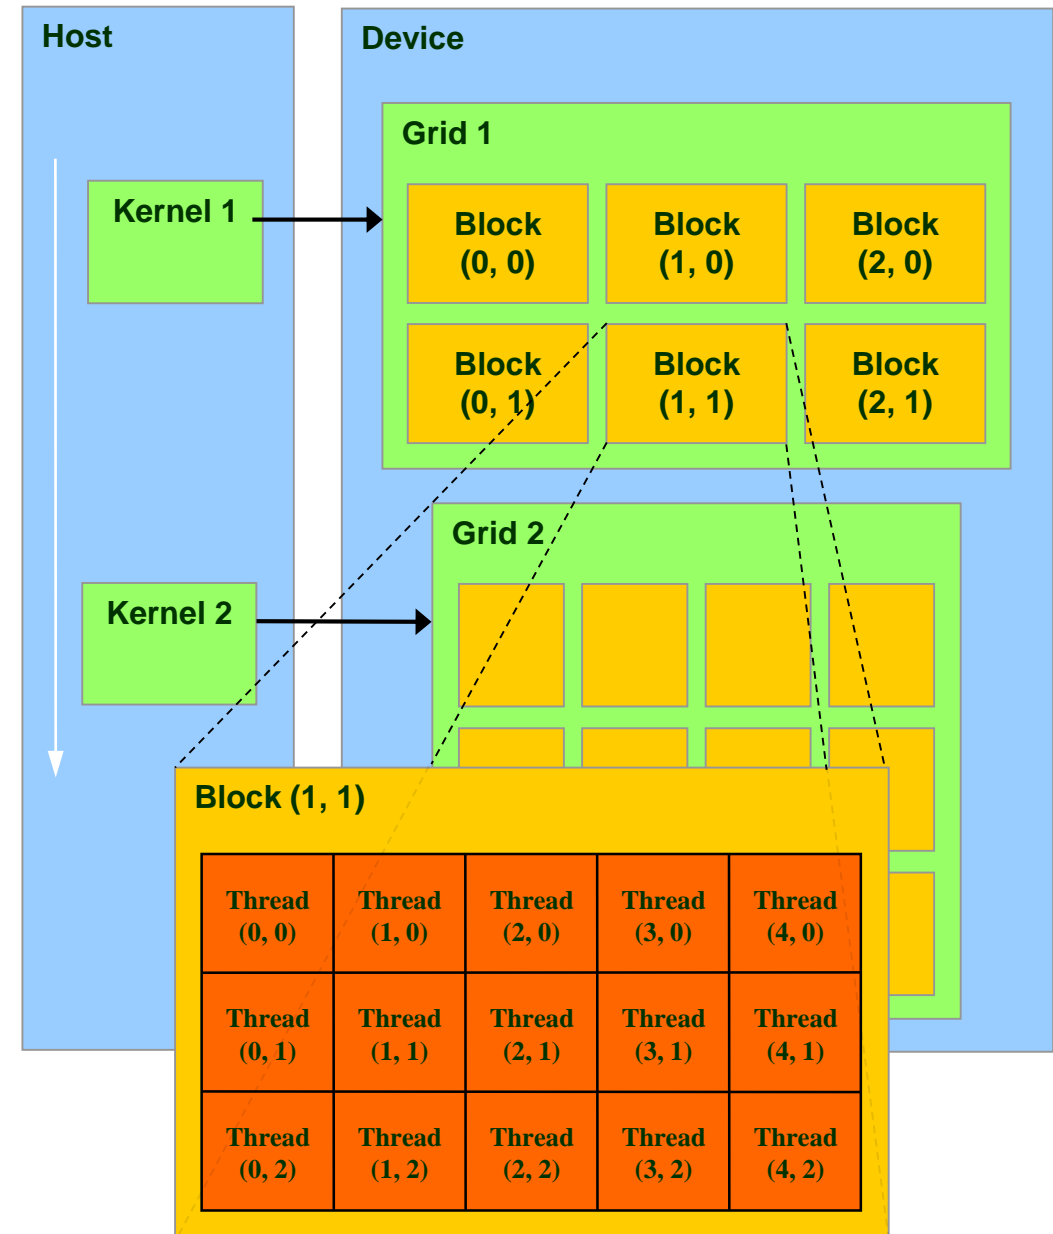
- Streaming Multiprocessor (SM)
  - 32 Streaming Processors (SP)
  - 4 Super Function Units (SFU)
- Multi-threaded instruction dispatch
  - 1 to **512 (1024, 2048)** threads active
  - Shared instruction fetch per 32 threads
  - Cover latency of texture/memory loads
- **64KB** shared memory / L1 cache
- DRAM texture and memory access





# Thread Life

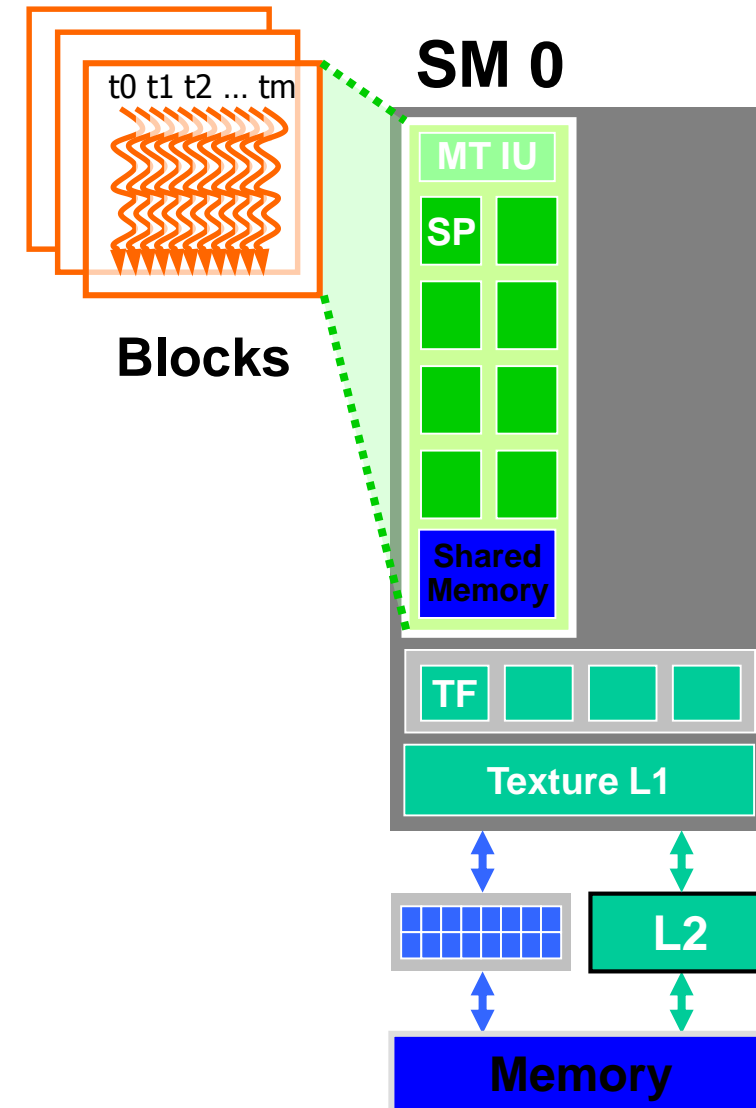
- Grid is launched on the SPA
- Thread Blocks are **serially** distributed to all the SM's
  - Potentially > 1 Thread Block per SM
- Each SM launches **Warps** of Threads
  - 2 levels of parallelism
- SM schedules and executes **Warps** that are ready to run
- As Warps and Thread Blocks complete, resources are freed
  - SPA can distribute more Thread Blocks





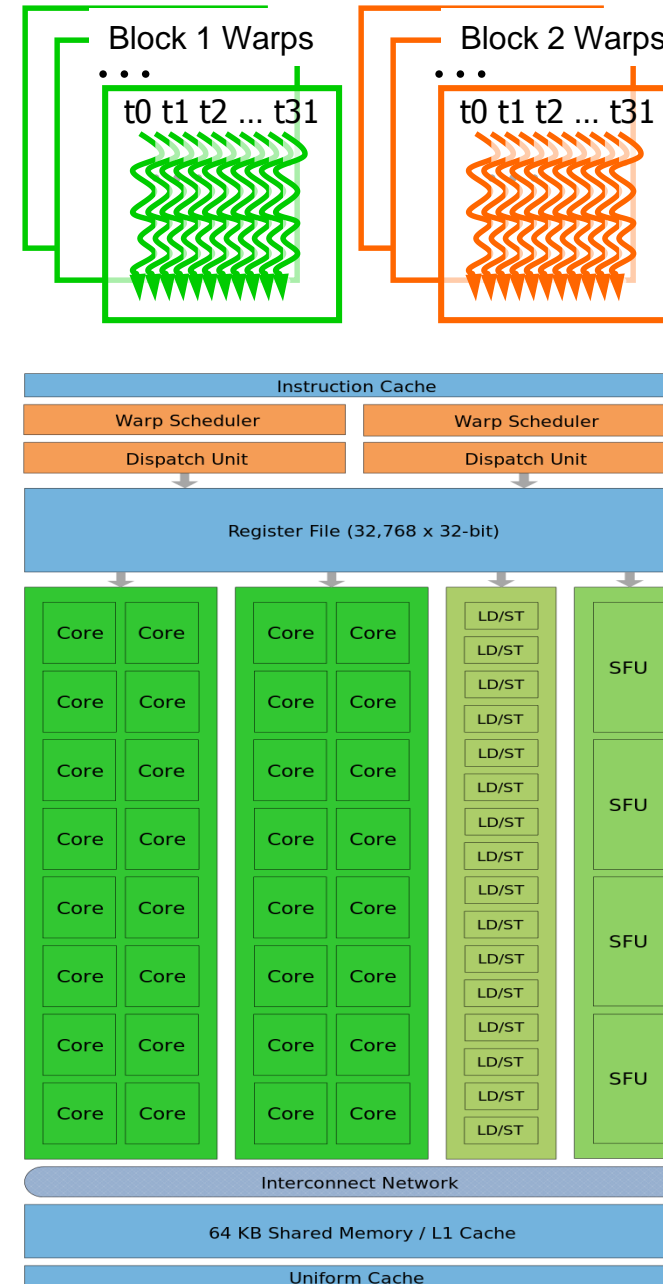
# Stream Multiprocessors Execute Blocks

- Threads are assigned to SMs in Block granularity
  - Up to **8 (and more)** Blocks to each SM as resource allows
  - SM in G200 can take up to 1K threads
    - Could be 256 (threads/block) \* 4 blocks
    - Or 128 (threads/block) \* 8 blocks, etc.
- Threads run concurrently
  - SM assigns/maintains thread id #s
  - SM manages/schedules thread execution

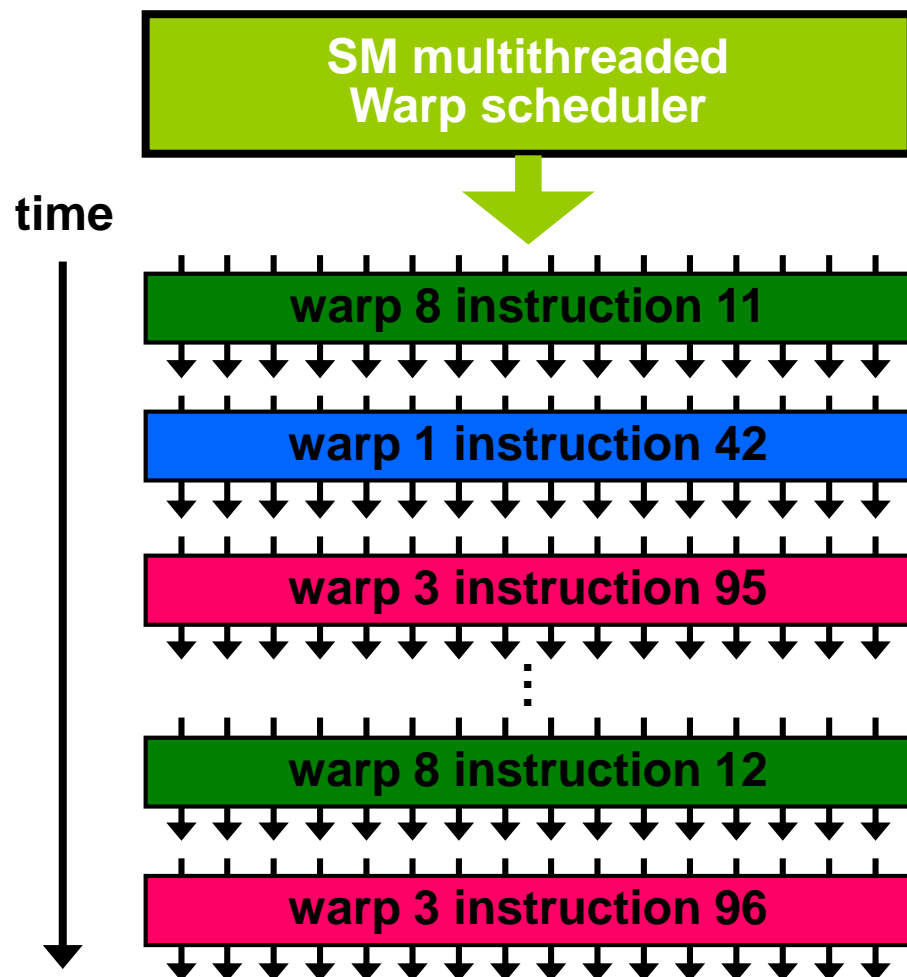


# Thread Scheduling and Execution

- Each Thread Blocks is divided in 32-thread Warps
  - This is an implementation decision, not part of the CUDA programming model
- Warp: primitive scheduling unit
- All threads in warp:
  - same instruction
  - control flow causes some to become inactive



# Warp Scheduling




- SM hardware implements zero-overhead Warp scheduling
  - Warps whose next instruction has its operands ready for consumption are eligible for execution
  - Eligible Warps are selected for execution on a prioritized scheduling policy
  - All threads in a Warp execute the same instruction when selected

## Memory System Goals

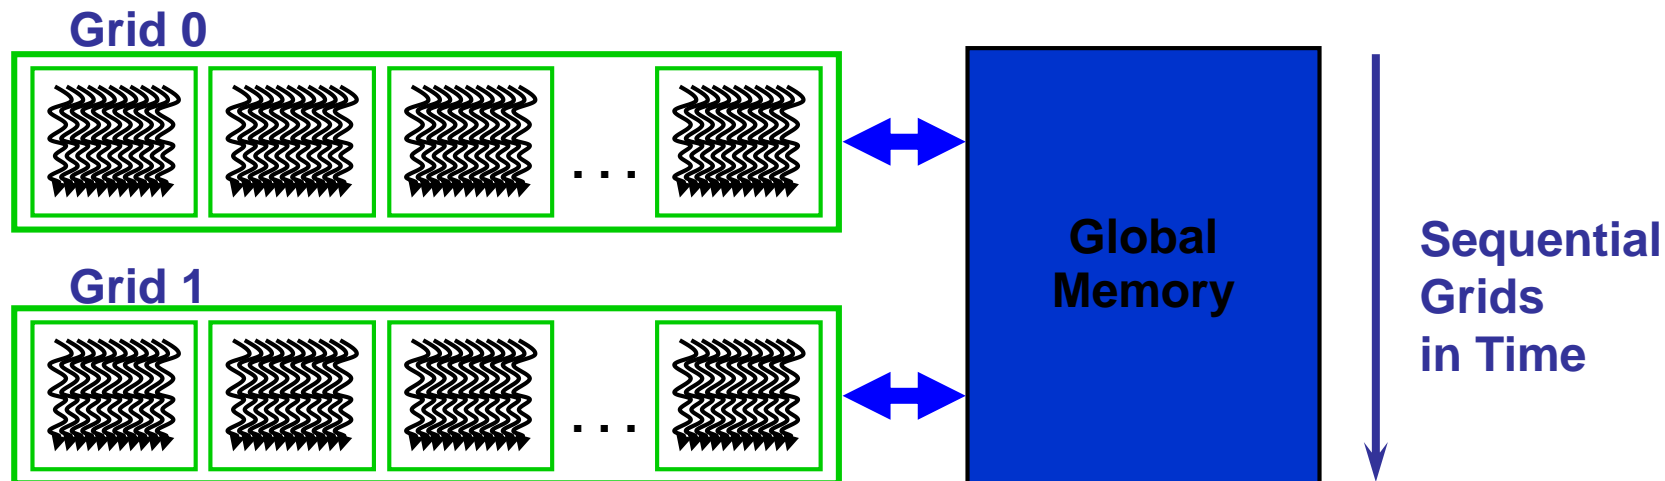
- High-Bandwidth (730 GB/s, 900 GB/s)
- As much parallelism as possible
- wide. 512 pins in G200 / Many DRAM chips
- max data rate per pin.
- maximize utilization
  - Multiple bins of memory requests
  - Coalesce requests to get as wide as possible
  - Goal to use every cycle to transfer from/to memory
- Caches where it makes sense. Small (L1, L2)



# Parallelism in the Memory System

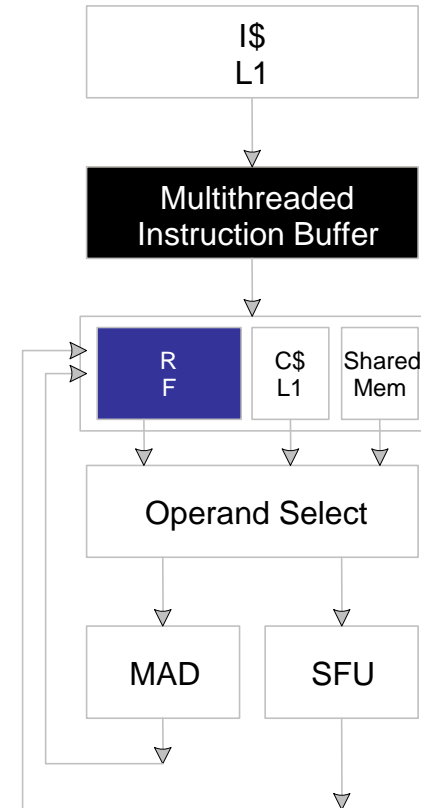


- Local Memory: per-thread
  - Private per thread
  - Auto variables, register
- Shared Memory: per-Block
  - Shared by threads of the same block
  - Inter-thread communication
- Global Memory: per-application
  - Shared by all threads
  - Inter-Grid communication



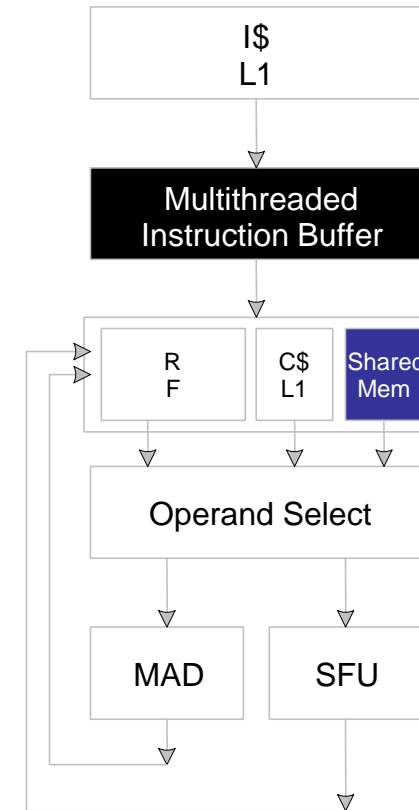
- Threads in a Block share data & results
  - In Memory and Shared Memory
  - Synchronize at barrier instruction
- Per-Block Shared Memory Allocation
  - Keeps data close to processor
  - Minimize trips to global Memory
  - SM Shared Memory dynamically allocated to Blocks, one of the limiting resources

- Register File (RF)
  - 64 KB (depending on arch)
  - 16K 32-bit registers
  - Provides 4 operands/clock
- TEX pipe can also read/write RF
  - 3 SMs share 1 TEX
- Load/Store pipe can also read/write RF



# Shared Memory


- Each SM has **64 (?)** KB of Shared Memory
  - **16** banks of 32bit words
- CUDA uses Shared Memory as shared storage visible to all threads in a thread block
  - read and write access
- Key Performance Enhancement
- Move data in Shared memory
- Operate in there



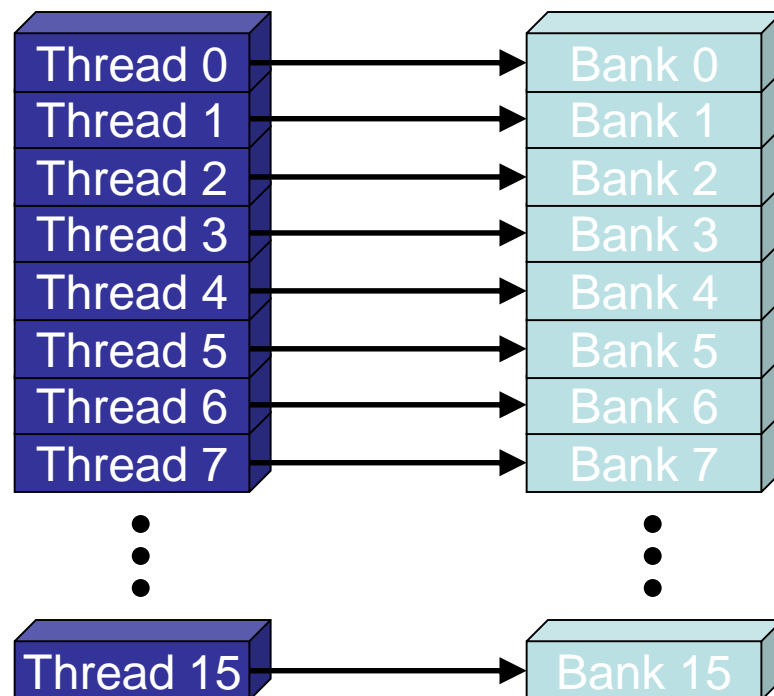


# Parallel Memory Architecture


- In a parallel machine, many threads access memory
  - Therefore, memory is divided into banks
  - Essential to achieve high bandwidth
- Each bank can service one address per cycle
  - A memory can service as many simultaneous accesses as it has banks
- Multiple simultaneous accesses to a bank result in a bank conflict
  - Conflicting accesses are serialized



- No Bank Conflicts
  - Linear addressing stride == 1

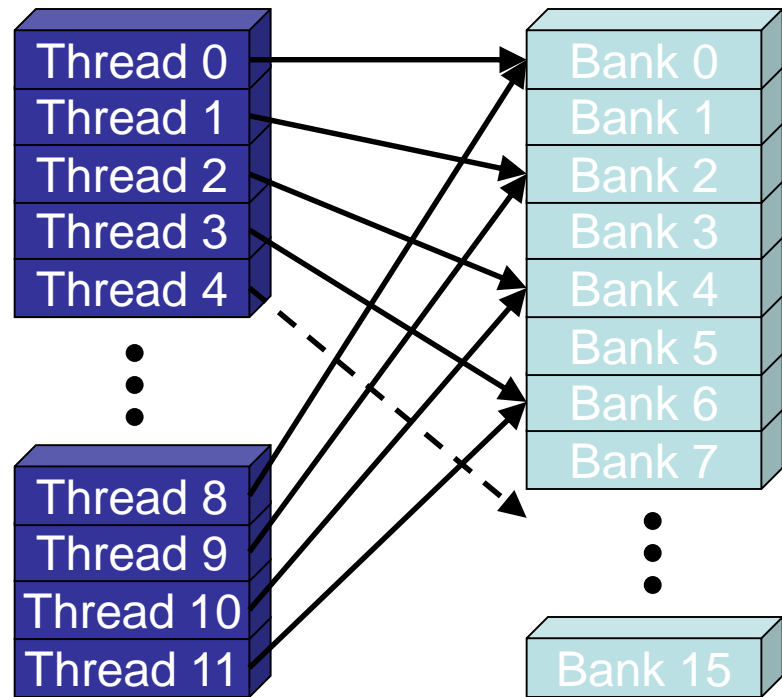


- No Bank Conflicts
  - Random 1:1 Permutation




- 2-way Bank Conflicts

- Linear addressing stride == 2



- 8-way Bank Conflicts

- Linear addressing stride == 8




# Shared memory bank conflicts


- Shared memory is as fast as registers if there are no bank conflicts
- The fast case:
  - If all threads of a half-warp access different banks, there is no bank conflict
  - If all threads of a half-warp access the identical address, there is no bank conflict (broadcast)
- The slow case:
  - Bank Conflict: multiple threads in the same half-warp access the same bank
  - Must serialize the accesses (or broadcast)
  - Cost = max # of simultaneous accesses to a single bank



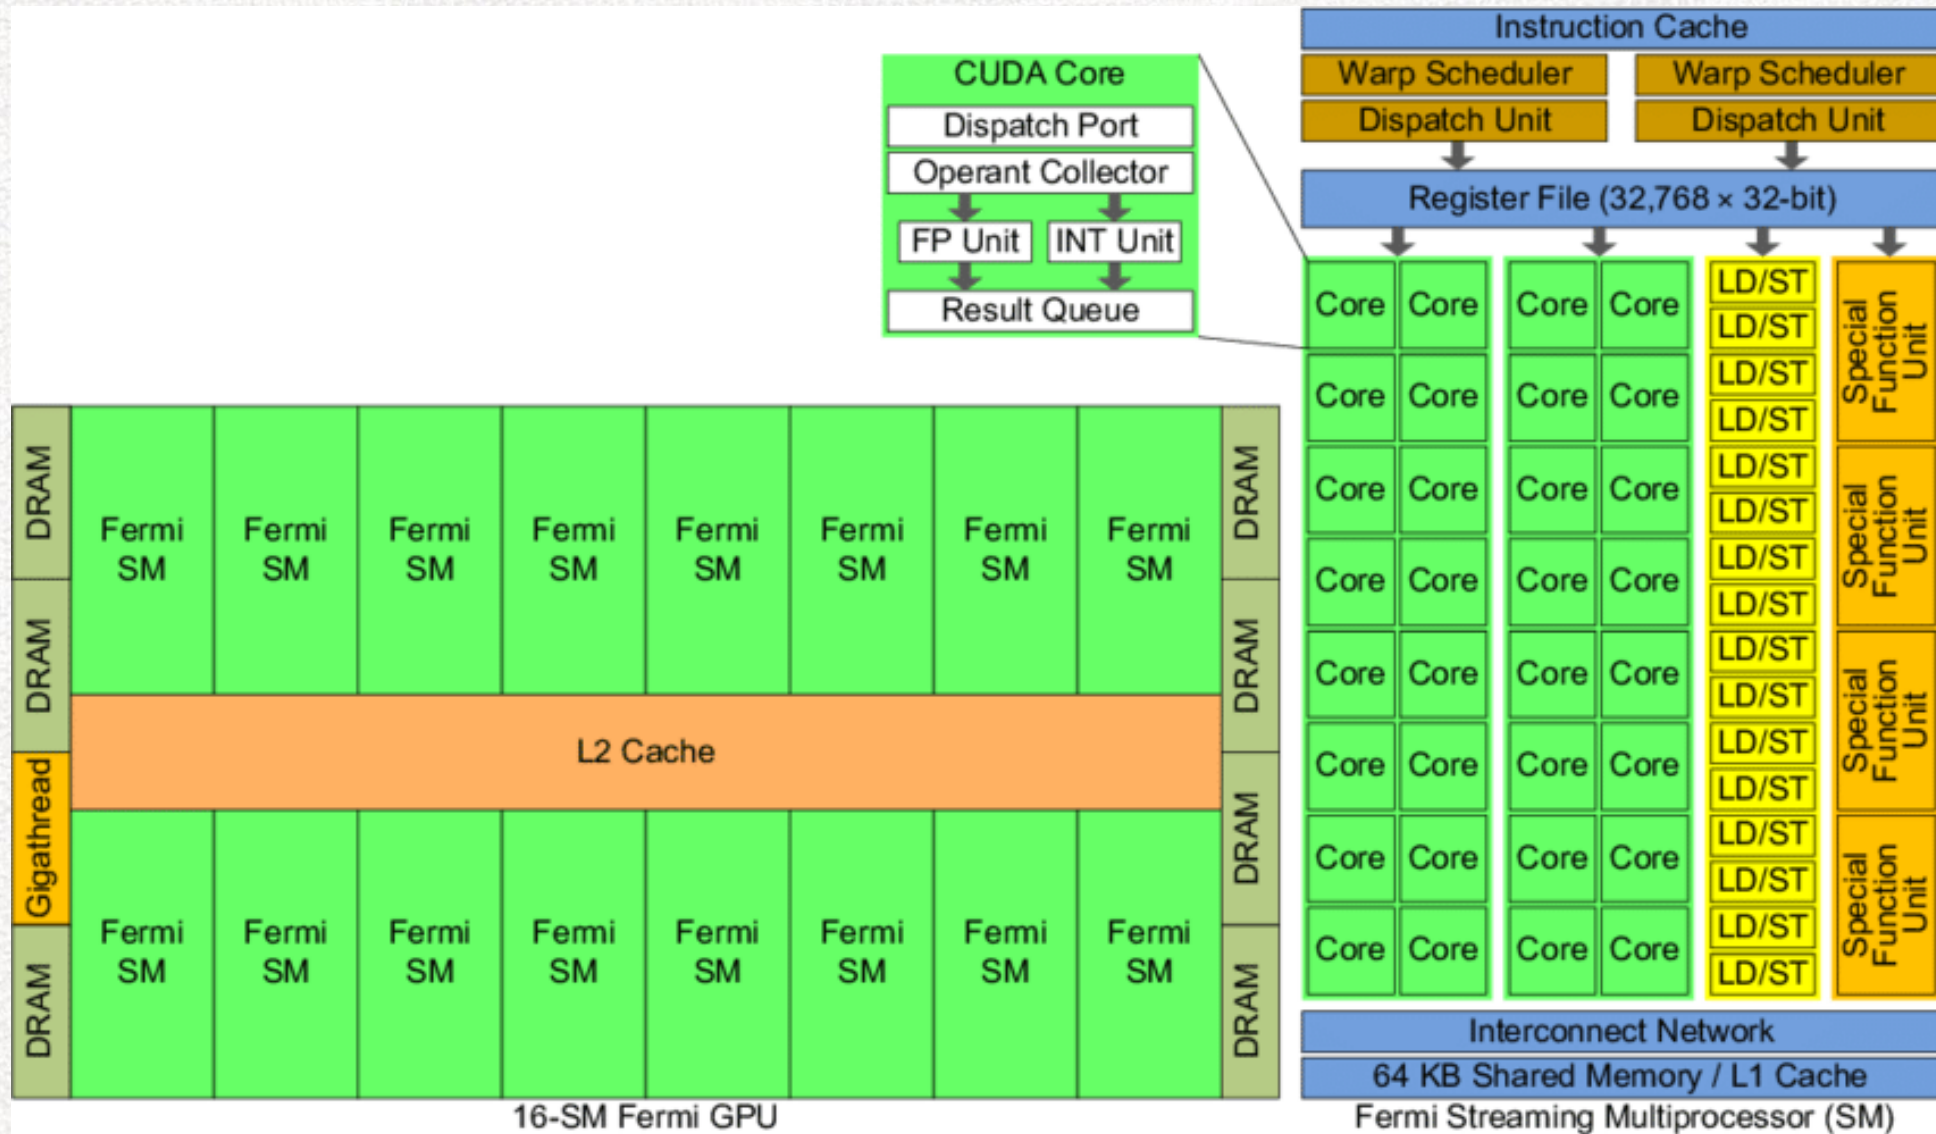
# Vector Reduction with Bank Conflicts



# No Bank Conflicts



# GPU 架构



**THANK YOU**

