

GPU Computing




规约算法

Hui Liu

Email: hui.sc.liu@gmail.com


Reduction Operations

- Multiple values are reduced into a single value
 - ADD, MUL, AND, OR,



- Useful primitive
- Easy enough to allow us to focus on optimization techniques

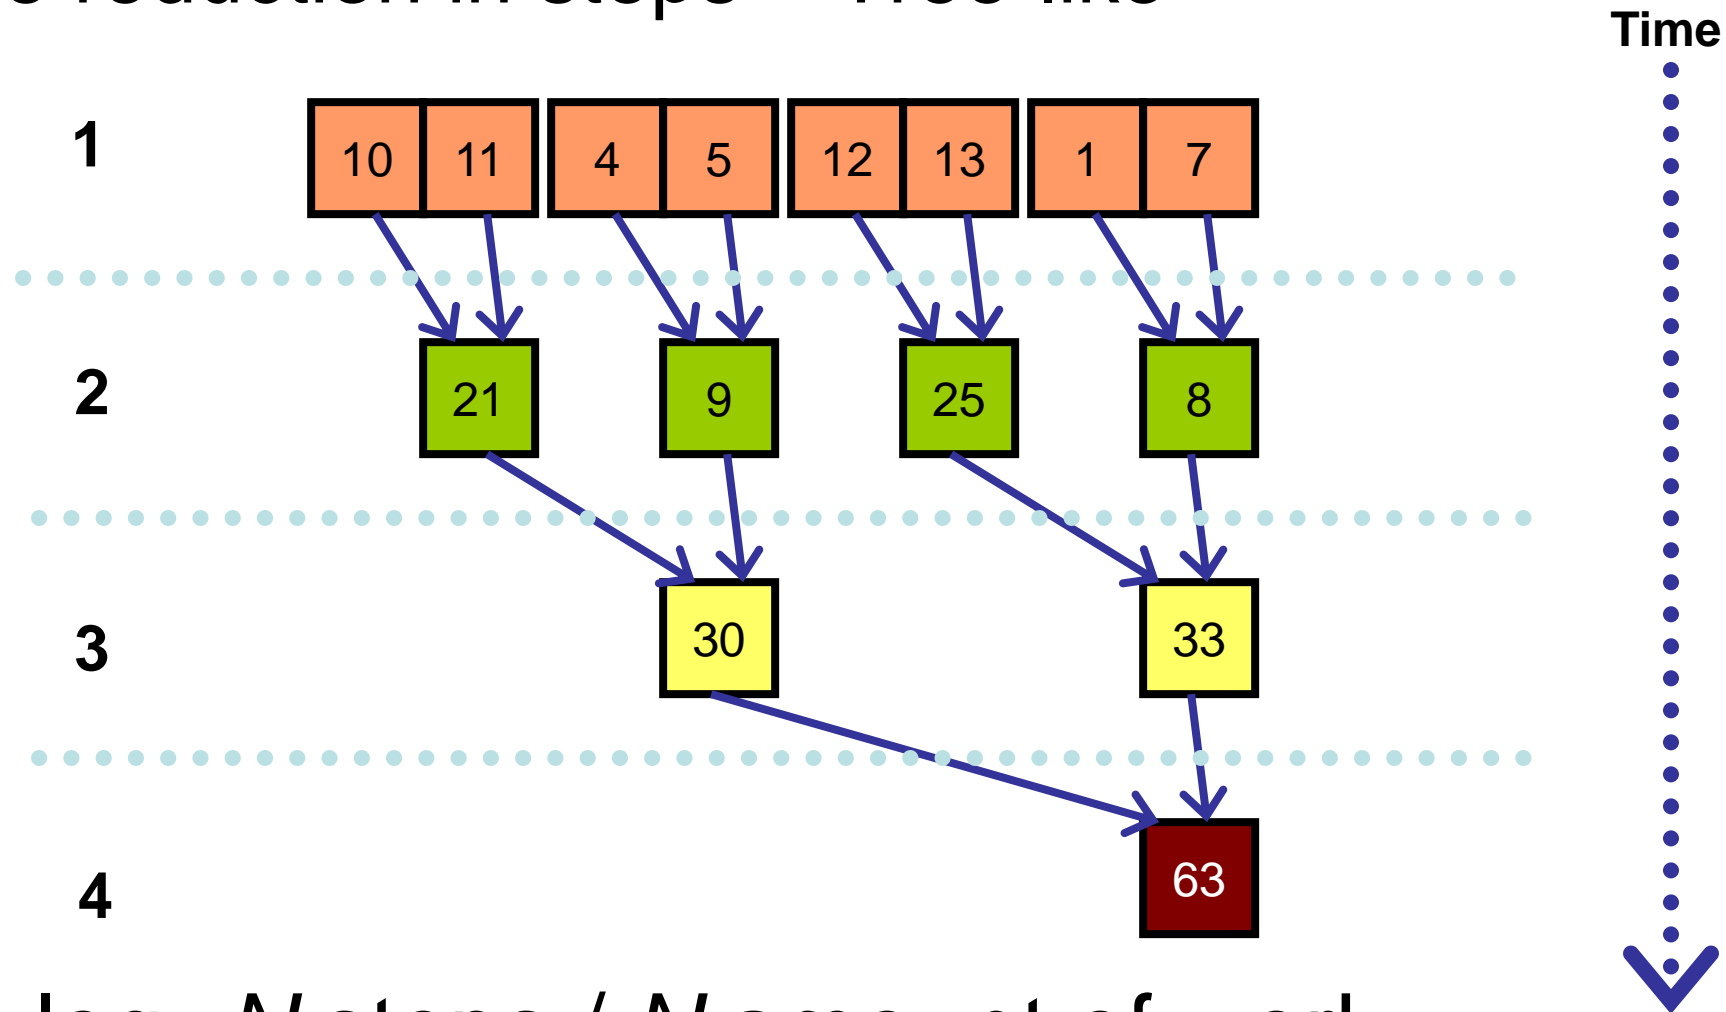
Sequential Reduction



- Start with the first two elements --> partial result
- Process the next element
- $O(N)$

Parallel Reduction

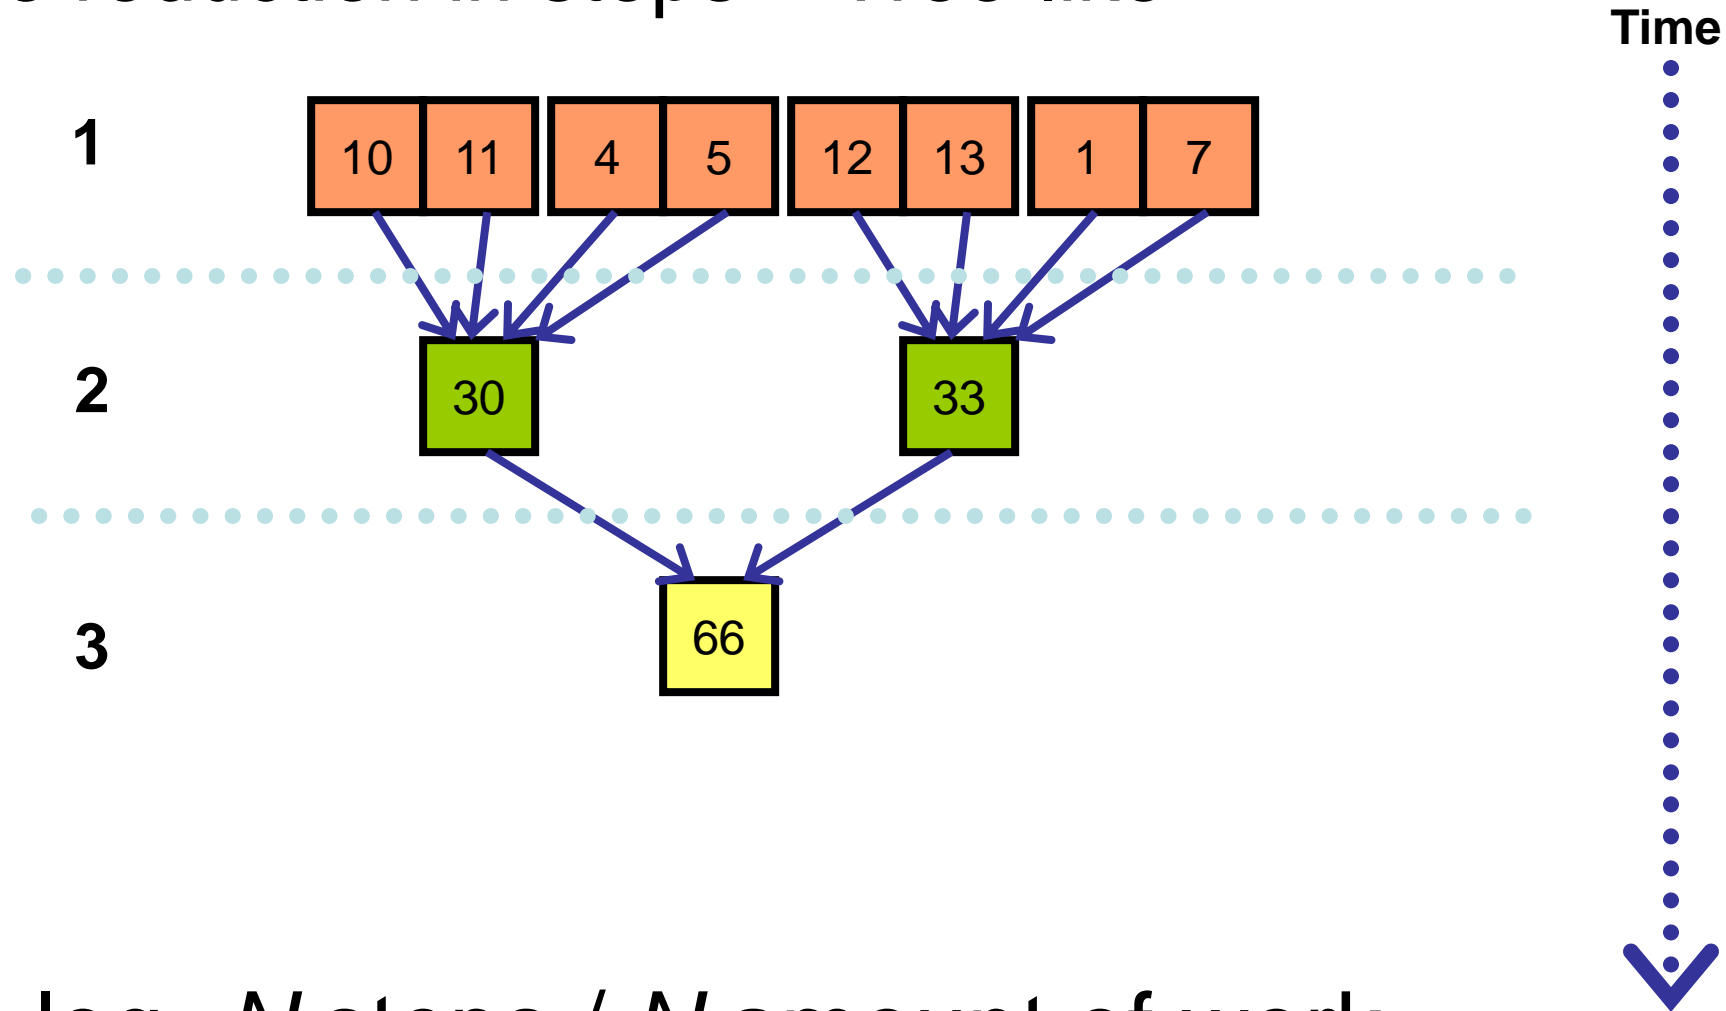
- Pair-wise reduction in steps – Tree-like



- $\log_2 N$ steps / N amount of work

Parallel Reduction – Different Degree-Trees possible


- Pair-wise reduction in steps – Tree-like



- $\log_4 N$ steps / N amount of work


CUDA Strategy

- Single Block:
 - Use Tree-Like approach
- Multiple Blocks?
 - Not a necessity
 - one thread can always process many elements
 - But, will suffer from low utilization
 - Utilize GPU resources
 - Useful for large arrays
- Each block processes a portion of the input



How about multiple blocks

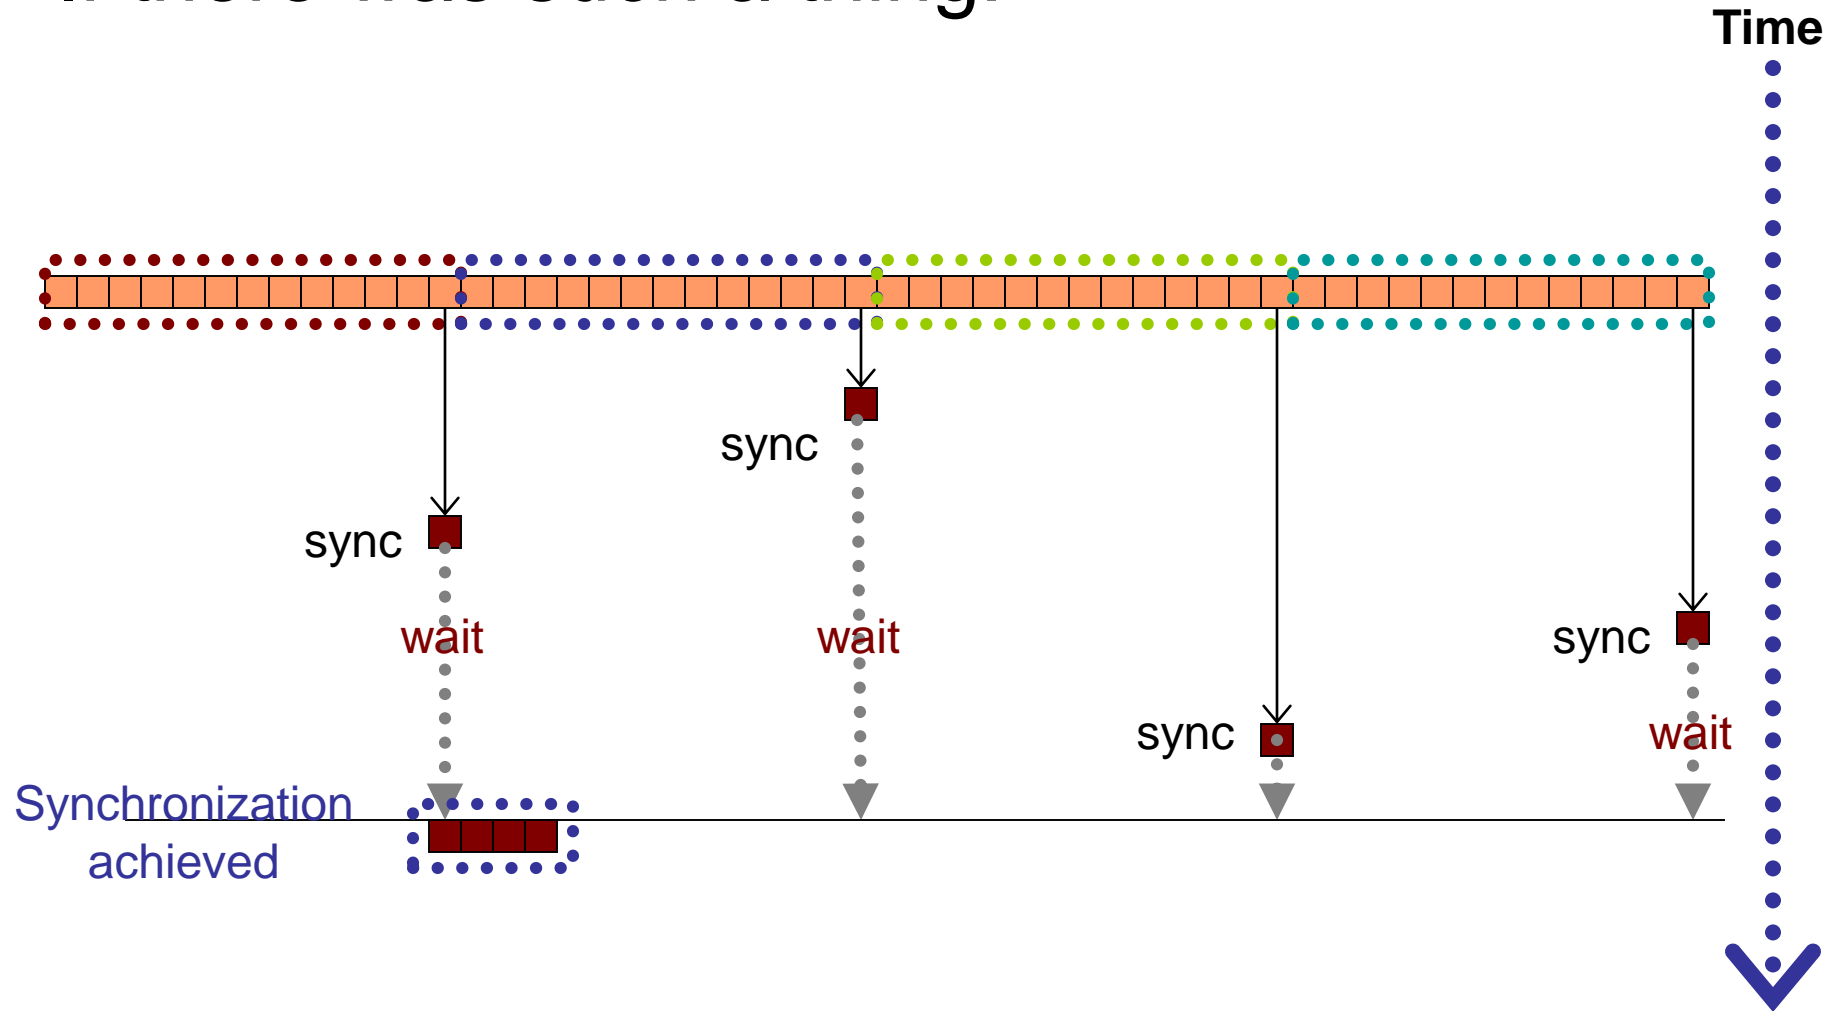
- How do we communicate results across blocks?



- The key problem is **synchronization**:
 - How do we know that each block has finished?


Global Synchronization

- If there was such a thing:

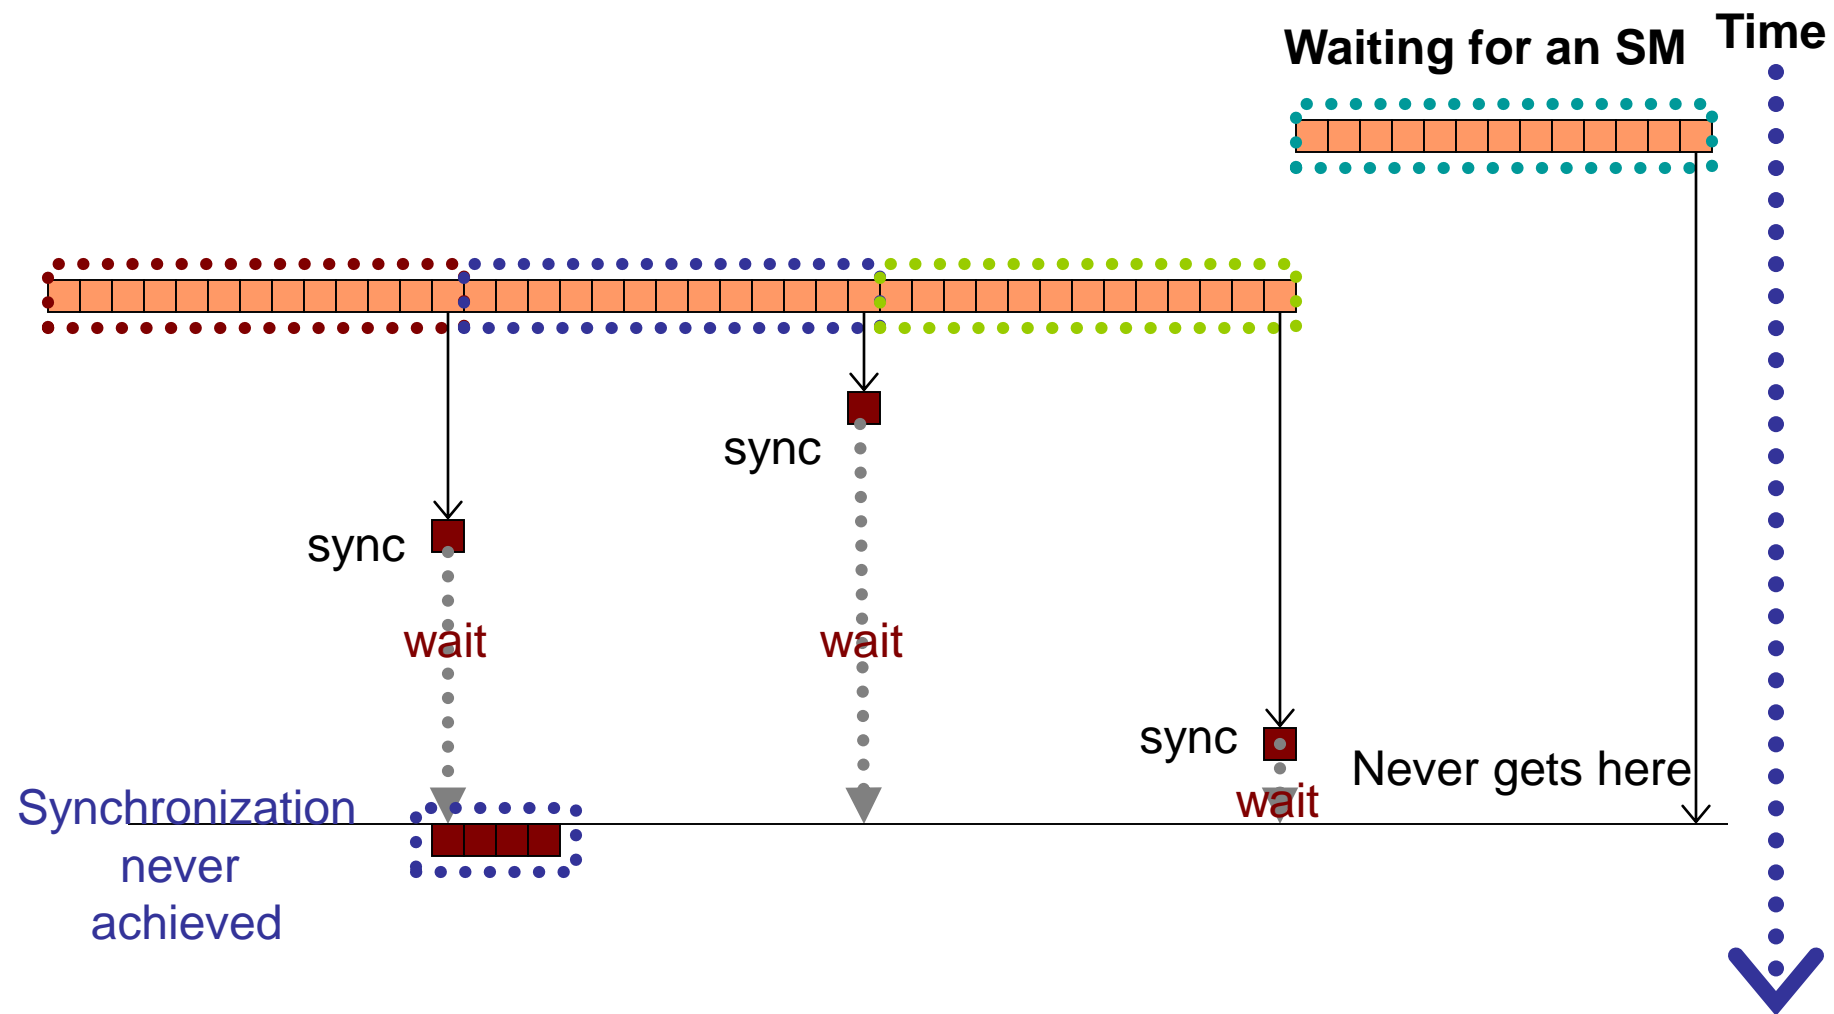


The Problem with Global Synchronization

- CUDA does not support it
 - One reason:
 - it's expensive to implement
 - Another reason:
 - “choose” between **limited blocks** or **deadlock**



Deadlock

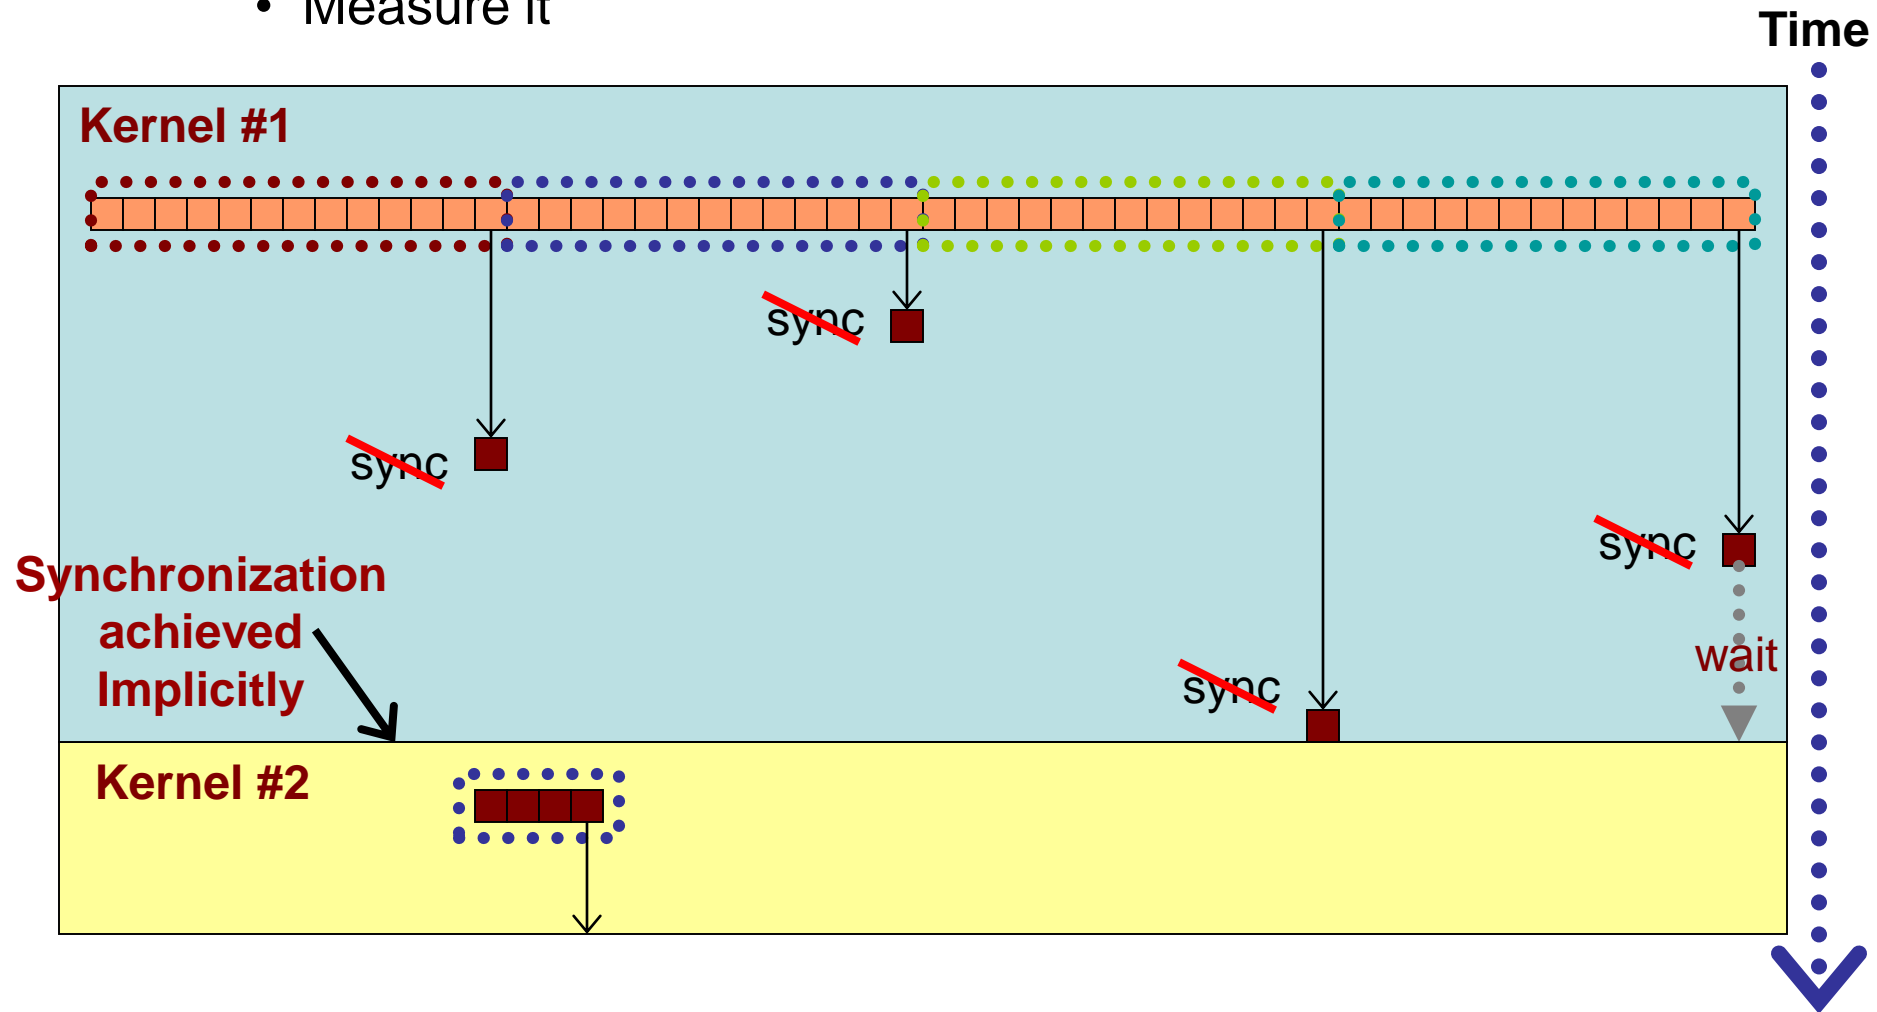


The Problem with Global Synchronization / Summary


- If there was global sync
 - Global sync after each block
 - Once all blocks are done, continue recursively
- CUDA does not have global sync:
 - Expensive to support
 - Would limit the number of blocks
 - Otherwise deadlock will occur
 - Once a block gets assigned to an SM it stays there
 - Each SM can take only 8 blocks
 - So, at most $\text{\#SMs} \times 8$ blocks could be active at any given point of time
- Solution: **Decompose into multiple kernels**

Decomposing into Multiple Kernels

- Implicit Synchronization between kernel invocations
 - Overhead of launching a new kernel non-negligible
 - Don't know how much
 - Measure it



Reduction: Big Picture



- The code for all levels is the same
- The same kernel code can be called multiple times

THANK YOU

