

GPU Computing



规约算法

Hui Liu

Email: hui.sc.liu@gmail.com

Reduction #6: Complete Unrolling

- If we knew the number of iterations at compile time, we could completely unroll the reduction
 - Block size is limited to 512
 - We can restrict our attention to powers-of-two block sizes
- We can easily unroll for a fixed block size
 - But we need to be generic
 - How can we unroll for block sizes we don't know at compile time?
- C++ Templates
 - CUDA supports C++ templates on device and host functions

Unrolling Using Templates

- Specify block size as a function template parameter:

```
template <unsigned int blockSize>  
__global__ void reduce5(int *g_data, int *g_odata)
```

Reduction #6: Completely Unrolled

```
if (blockSize >= 512) {  
    if (tid < 256) { sdata[tid] += sdata[tid + 256]; } __syncthreads();  
}  
if (blockSize >= 256) {  
    if (tid < 128) { sdata[tid] += sdata[tid + 128]; } __syncthreads();  
}  
if (blockSize >= 128) {  
    if (tid < 64) { sdata[tid] += sdata[tid + 64]; } __syncthreads();  
}  
if (tid < 32) {  
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];  
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];  
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];  
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];  
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];  
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];  
}
```

- Note: all code in **RED** will be evaluated at compile time.
- Results in a very efficient inner loop

What if block size is not known at compile time?

- There are “only” 10 possibilities:

switch (threads)

{

case 512:

reduce5<512><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 256:

reduce5<256><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 128:

reduce5<128><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 64:

reduce5< 64><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 32:

reduce5< 32><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 16:

reduce5< 16><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 8:

reduce5< 8><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 4:

reduce5< 4><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 2:

reduce5< 2><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

case 1:

reduce5< 1><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;

}

Performance for 4M element reduction

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing non-divergent branching	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first step during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: Unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
Kernel 6: Complete Unroll	0.381 ms	43.996 GB/s	1.41x	21.16x

THANK YOU

