

GPU Computing



CUDA 程序优化

Hui Liu

Email: hui.sc.liu@gmail.com

优化内存使用

基本策略

- 计算 vs 内存访问
 - Processing data is cheaper than moving it around
 - Especially for GPUs as they devote many more transistors to ALUs than memory
- And will be increasingly so
 - The less memory bound a kernel is, the better it will scale with future GPUs
- So you want to:
 - Maximize use of low-latency, high-bandwidth memory
 - Optimize memory access patterns to maximize bandwidth
 - Leverage parallelism to hide memory latency by overlapping memory accesses with computation as much as possible
 - Kernels with high arithmetic intensity (ratio of math to memory transactions)
 - Sometimes recompute data rather than cache it

极小化 CPU \leftrightarrow GPU 数据传输

- CPU \leftrightarrow GPU memory bandwidth much lower than GPU memory bandwidth
 - Use page-locked host memory (`cudaMallocHost()`) for maximum CPU \leftrightarrow GPU bandwidth
 - Be cautious however since allocating too much page-locked memory can reduce overall system performance
- Minimize CPU \leftrightarrow GPU data transfers by moving more code from CPU to GPU
 - Even if that means running kernels with low parallelism computations
 - Intermediate data structures can be allocated, operated on, and deallocated without ever copying them to CPU memory
- Group data transfers
 - One large transfer much better than many small ones

极大化使用共享内存

- Shared memory is hundreds of times faster than global memory
- Threads can cooperate via shared memory
 - Not so via global memory
- Common CUDA kernel structure:
 1. Load data from global memory to shared memory
 2. `__syncthreads()`
 3. Process the data in shared memory with many threads
 4. `__syncthreads()` (if needed)
 5. Store results from shared memory to global memory
- Notes:
 - Steps 2-4 may be repeated, looped, etc.
 - Step 4 is not necessary if there is no dependence of stored data on other threads


优化内存使用模式

- Effective bandwidth can vary by an order of magnitude depending on access pattern
- Optimize access patterns to get:
 - *Coalesced* global memory accesses
 - Shared memory accesses with *no or few bank conflicts*
 - *Cache-efficient* texture memory accesses
 - *Same-address* constant memory accesses


全局内存: 对齐与合并访问

- The simultaneous global memory accesses by each thread of a half-warp during the execution of a single read or write instruction will be *coalesced* into a single access if:
 - The size of the memory element accessed by each thread is either 4, 8, or 16 bytes
 - The elements form a contiguous block of memory
 - The N^{th} element is accessed by the N^{th} thread in the half-warp
 - The address of the first element is aligned to 16 times the element's size
- Coalescing happens even if some threads do not access memory (divergent warp)

全局内存合并访问




Coalesced `float` memory access




Coalesced `float` memory access

(divergent warp)

非合并全局内存访问




Non-sequential `float` memory access




Misaligned starting address

非合并全局内存访问




Non-contiguous `float` memory access



Non-coalesced `float3` memory access

共享内存

- In a parallel machine, many threads access memory
 - Therefore, memory is divided into *banks*
 - Essential to achieve high bandwidth
- Each bank can service one address per cycle
 - A memory can service as many simultaneous accesses as it has banks
- Multiple simultaneous accesses to a bank result in a *bank conflict*
 - Conflicting accesses are serialized



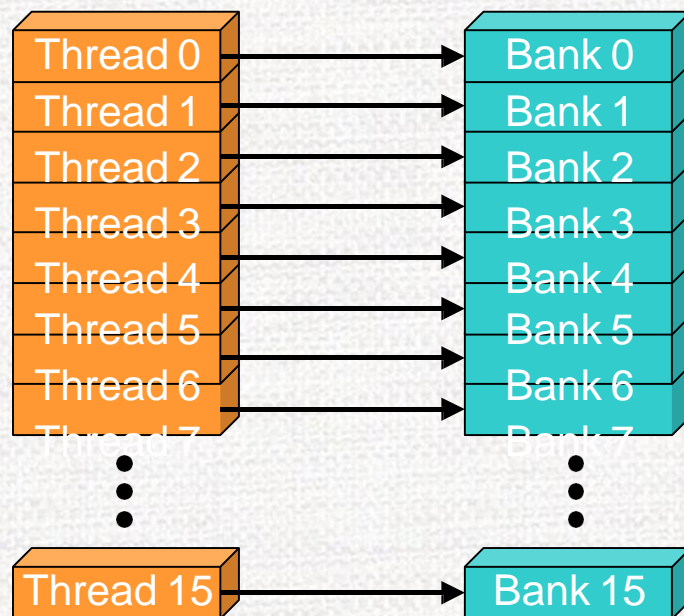
共享内存: banked

- Bandwidth of each bank is 32 bits per 2 clock cycles
- Successive 32-bit words are assigned to successive banks
- G80 has 16 banks
 - So bank = address % 16
 - Same as the size of a half-warp
 - No bank conflicts between different half-warps, only within a single half-warp

共享内存访问


- No bank conflicts

- Linear addressing
stride == 1



- No bank conflicts

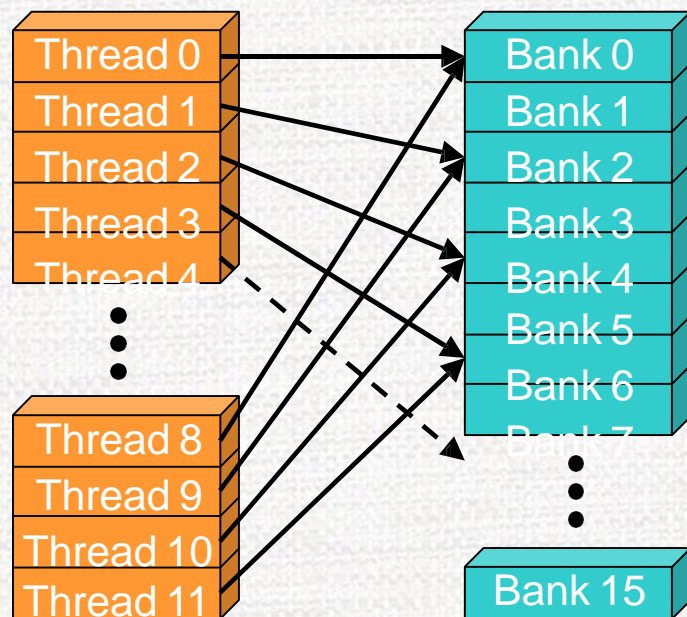
- Random 1:1 permutation



共享内存访问


- 2-way bank conflicts

- Linear addressing
stride == 2



- 8-way bank conflicts

- Linear addressing
stride == 8



共享内存访问冲突

- Shared memory is as fast as registers if there are no bank conflicts
- The fast case:
 - If all threads of a half-warp access different banks, there is no bank conflict
 - If all threads of a half-warp read the same word, there is no bank conflict (*broadcast*)
- The slow case:
 - Bank conflict: multiple threads in the same half-warp access the same bank
 - Must serialize the accesses
 - Cost = max # of simultaneous accesses to a single bank

优化优先级

- Coalesce global memory accesses
 - #1 priority, highest !/\$ optimization
- Take advantage of shared memory
- Hide memory latency by running many threads with high arithmetic intensity
- Leave bank conflicts for last!
 - 4-way and smaller conflicts are not usually worth avoiding if it will cost more instructions

THANK YOU

