

Texture Memory

Texture Memory

- read-only memory
- Can improve performance and reduce memory traffic when reads have certain access patterns.
- Originally designed for the classical OpenGL and DirectX rendering pipelines.
- But has some properties that make it extremely useful for computing, especially computer vision application.

Texture Memory

- Texture memory is cached on chip
 - In KB range in every SM
 - In some situations it will provide higher effective bandwidth by reducing memory requests to off-chip DRAM.
- Texture caches are designed for graphics applications where memory access patterns exhibit a great deal of **spatial locality**.
- Every SM has several texture fetch units




One SM from Pascal architecture (GP100)

Source: <https://www.extremetech.com/extreme/226032-nvidias-pascal-gp100-gpu-massive-bandwidth-enormous-double-precision-performance>

Texture Memory

- The texture cache is optimized for 2D spatial locality.
- Part of DRAM
- The process of reading a texture is called a *texture fetch*.
- Can be addressed as 1D, 2D, or 3D dimensional arrays.
- Elements of the array are called *texels*.

texture references



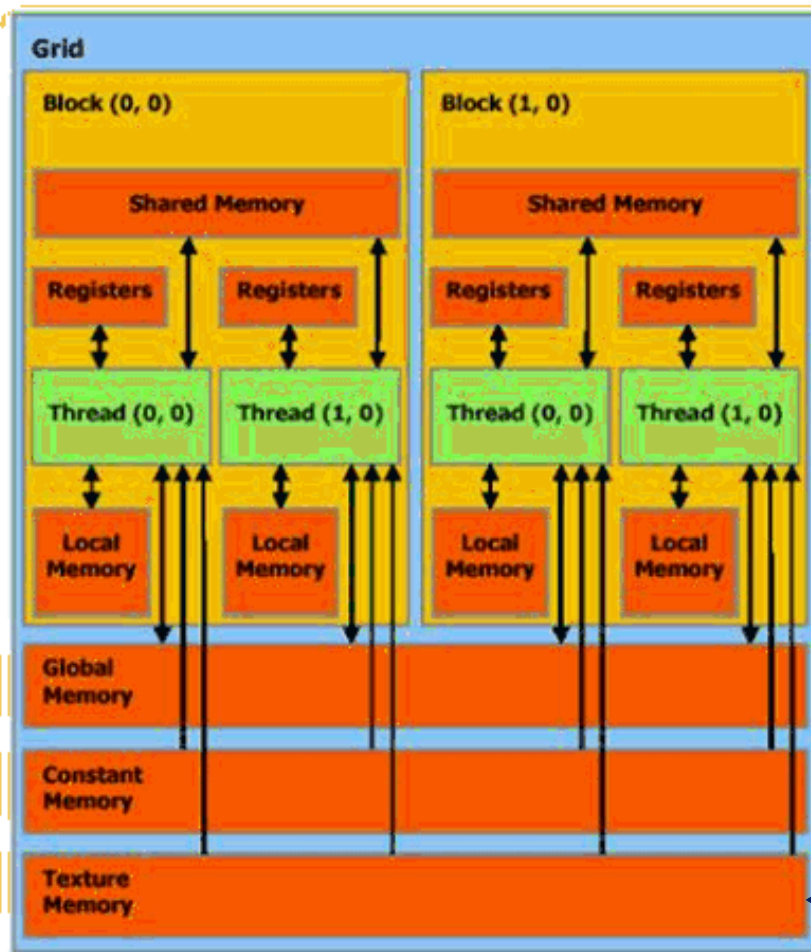
CUDA Arrays

- Designed specifically to support texturing.
- Allocated from device memory
- Do not consume any CUDA address space.
- Have an **opaque layout**
- Cannot be addressed by pointers
- Memory locations addressed using two things:
 - array handle
 - set of 1D, 2D, or 3D coordinates


Why CUDA Arrays?

- Designed so that contiguous addresses exhibit 2D or 3D locality.

Texture Memory


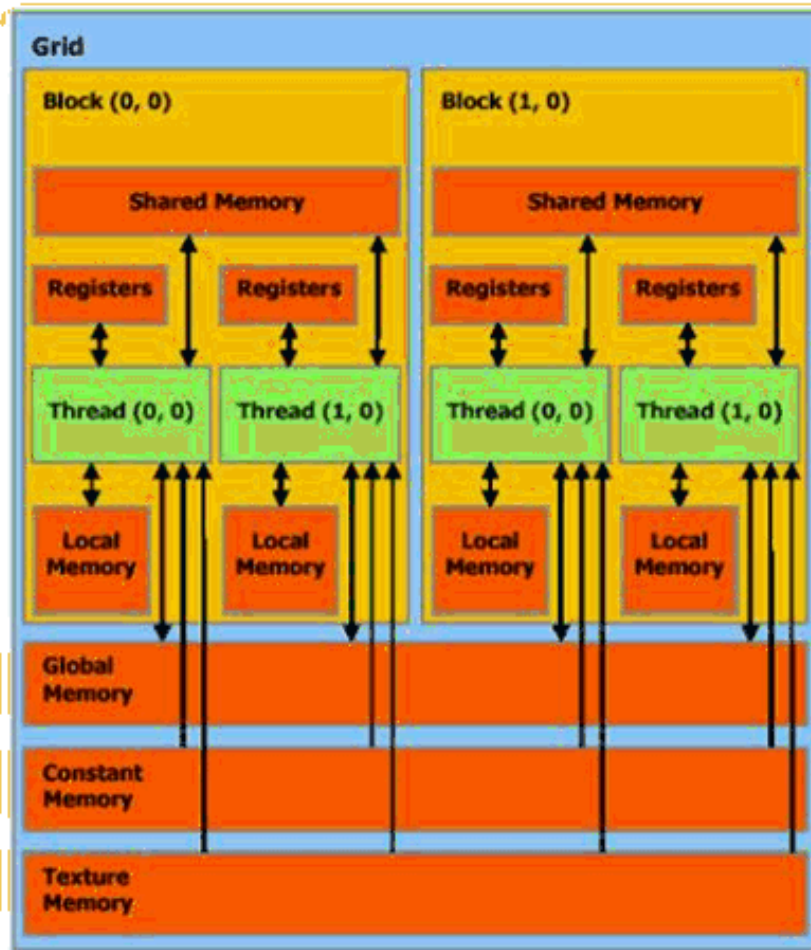


To accelerate frequently performed operations such as mapping a 2D "skin" onto a 3D polygonal model.

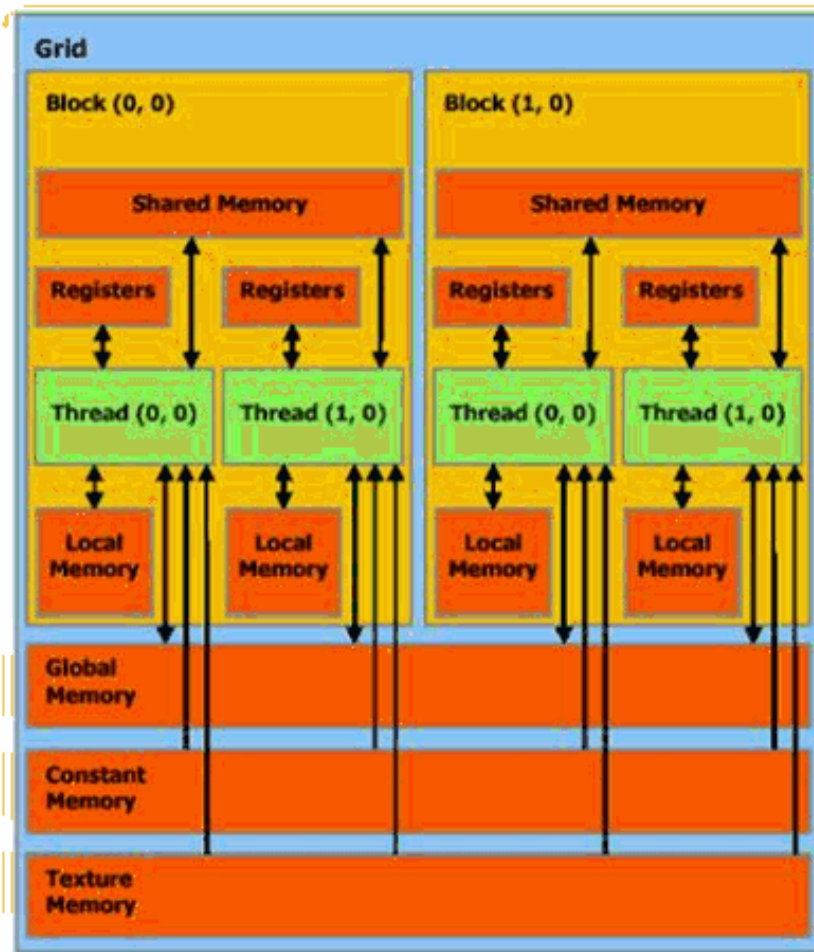


source: <http://cuda-programming.blogspot.com/2013/02/texture-memory-in-cuda-what-is-texture.html>

Texture Memory



Texture Memory



Capabilities:

- Ability to cache global memory
- Dedicated interpolation hardware
- Provides a way to interact with the display capabilities of the GPU.

The best performance is achieved when the threads of a warp read locations that are close together from a spatial locality perspective.

Allocating CUDA Arrays

```
cudaError_t cudaMallocArray (
    struct cudaArray **      array,
    const struct cudaChannelFormatDesc * desc,
    size_t    width,
    size_t    height,
    unsigned int flags
);
```

array: pointer to allocated array in device memory

width: array width in bytes

height: default is 0 → 1D array

flags:

- `cudaArrayDefault`: default array allocation
- `cudaArraySurfaceLoadStore`

Allocating CUDA Arrays

```
cudaError_t cudaMallocArray (  
    struct cudaArray ** array,  
    const struct cudaChannelFormatDesc * desc,  
    size_t    width,  
    size_t    height,  
    unsigned int flags  
);
```

struct cudaChannelFormatDesc
cudaCreateChannelDesc(x,y,z,w,f);

```
struct cudaChannelFormatDesc {  
    int x, y, z, w; ← number of bits in each member of the texture element  
    enum cudaChannelFormatKind f;  
};
```

cudaChannelFormatKindSigned
cudaChannelFormatKindUnsigned
cudaChannelFormatKindFloat
cudaChannelFormatKindNone

Signed channel format
Unsigned channel format
Float channel format
No channel format

Texture Fetch

- First parameter is texture reference
 - defines which part of texture memory is fetched
 - must be bound through runtime functions to texture memory
 - Attribute:
 - texture is addressed as 1D, 2D, or 3D
 - the input and output data types of the texture fetch
 - the input coordinates are interpreted
 - what processing should be done
 - Type of texels are the basic: integer, single/double precision floating point,

Steps for Using Texture Memory in Your CUDA Code

1. **Declare** the texture memory in CUDA.
2. **Bind** the texture memory to your texture reference in CUDA.
3. **Read** the texture memory from your texture reference in CUDA Kernel.
4. **Unbind** the texture memory from your texture reference in CUDA.

Step 1: Declare

texture <type, dim, readmode> texture_reference;

- texture_reference: the handle to be used
- type: type of texel data returned from an access to the texture: int, float,
- dim: 1 (default), 2, or 3
- readmode: controls conversion of texel returned by an access
 - cudaReadModeElementType (default) no conversion
 - cudaReadModeNormalizedFloat
 - if type is integer, value returned is mapped to [-1.0,1.0] for signed, and [0.0, 1.0] for unsigned
- Example:

texture <float, 2, cudaReadModeElementType> mytex;

Step 2: Bind

cudaBindtexture (size *t offset,
& texture_reference , const void * devptr,
size_t size) ;

- Binds size bytes of the memory area pointed to by devPtr to texture reference texture_reference.
- offset parameter is an optional byte offset.
- devPtr: Memory area on device
- size: Size of the memory area pointed to by devPtr

Step 3: Read

- The easiest is: **tex1Dfetch()**

Example:

```
texture<int,1,cudaReadModeElementType> texref;  
__global__  
void textureTest(int *out){  
    int tid = blockIdx.x * blockDim.x + threadIdx.x;  
    float x;  
    int i;  
    for(i=0; i<30; i++)  
        x = tex1Dfetch(texref, i);  
}
```

Step 4: Unbind

```
cudaUnbindTexture(texture_reference);
```

So

- Texture memory size is very small.
- We just scratched the surface of texture memory.
- Two usages of texture memory outside graphics applications:
 - Using texture cache to reduce bandwidth and work around coalescing constraints.
 - Make use of advanced fixed-function hardware put inside GPU for graphics applications